Pollution Image Classification

Artificial Vision - Project

Professor: Lanzarotti Raffaella
Student: Francesco Ruggeri - 871798

A.A. 2023/2024

Introduction

Objective

Develop a waste image classification model using artificial vision techniques

Dataset

WaRP
Waste Recycling Plant Dataset

- Sourced in Kaggle
- 28 distinct classes
- 10000 images

Dataset Exploration

Aspect-Ratio Distribution

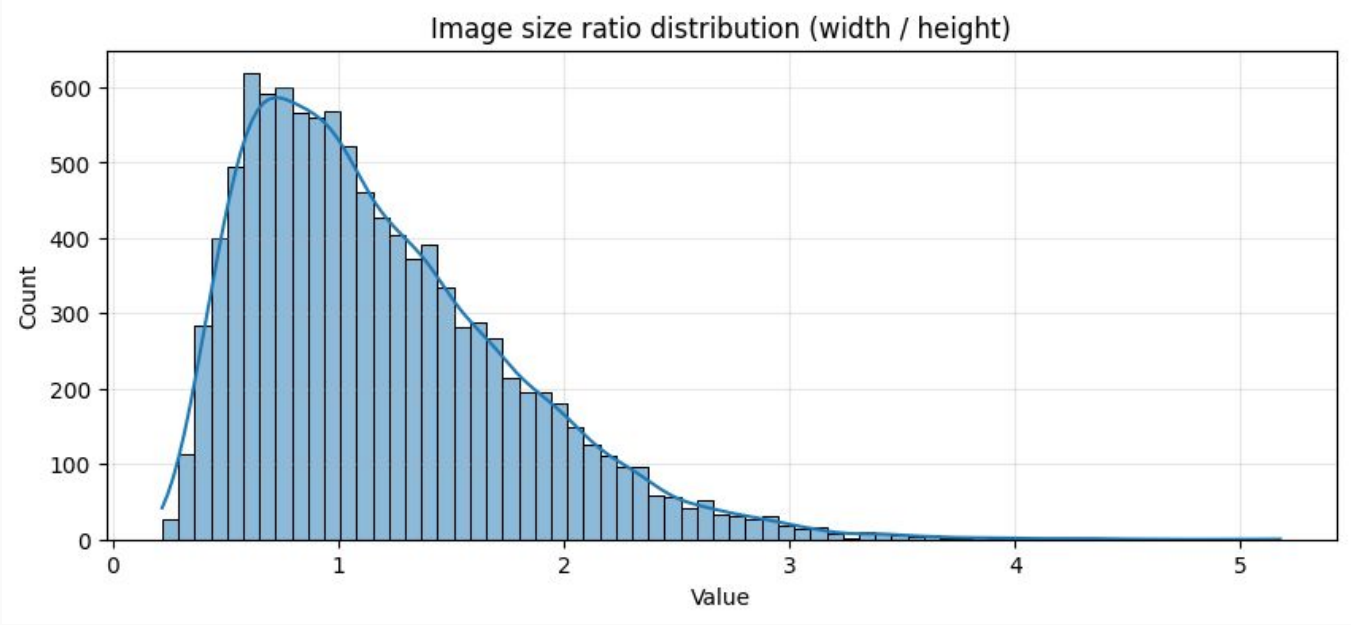

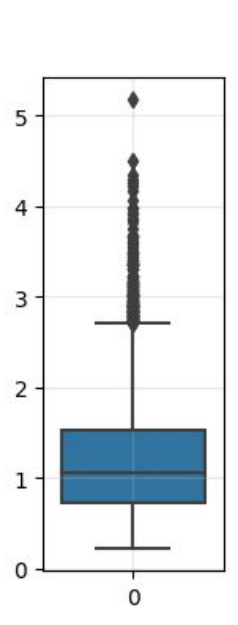

Dataset Exploration

Images Width and Height Distribution

Optimal width, height: 128, 128

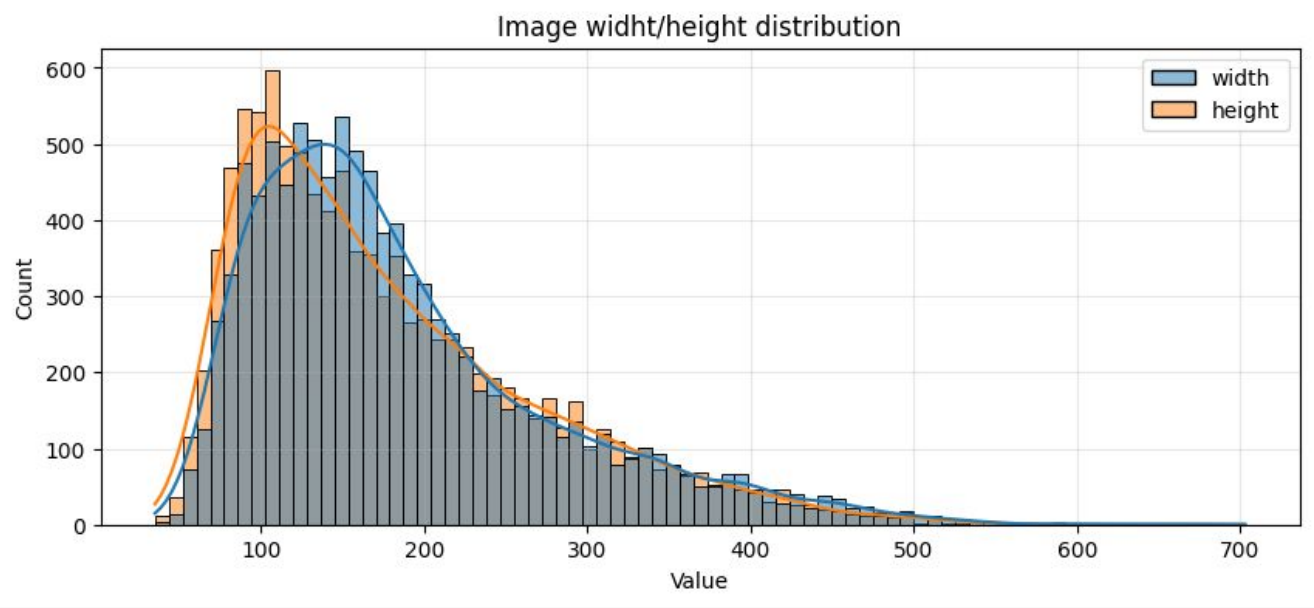

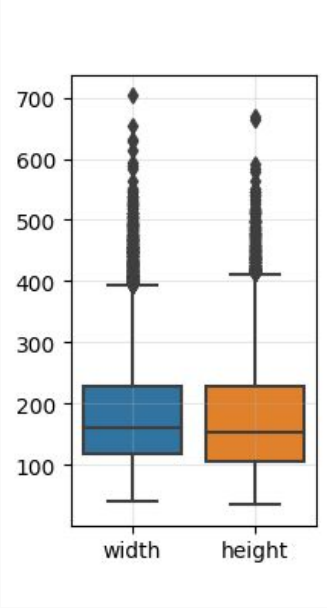

Dataset Exploration

Class Distribution

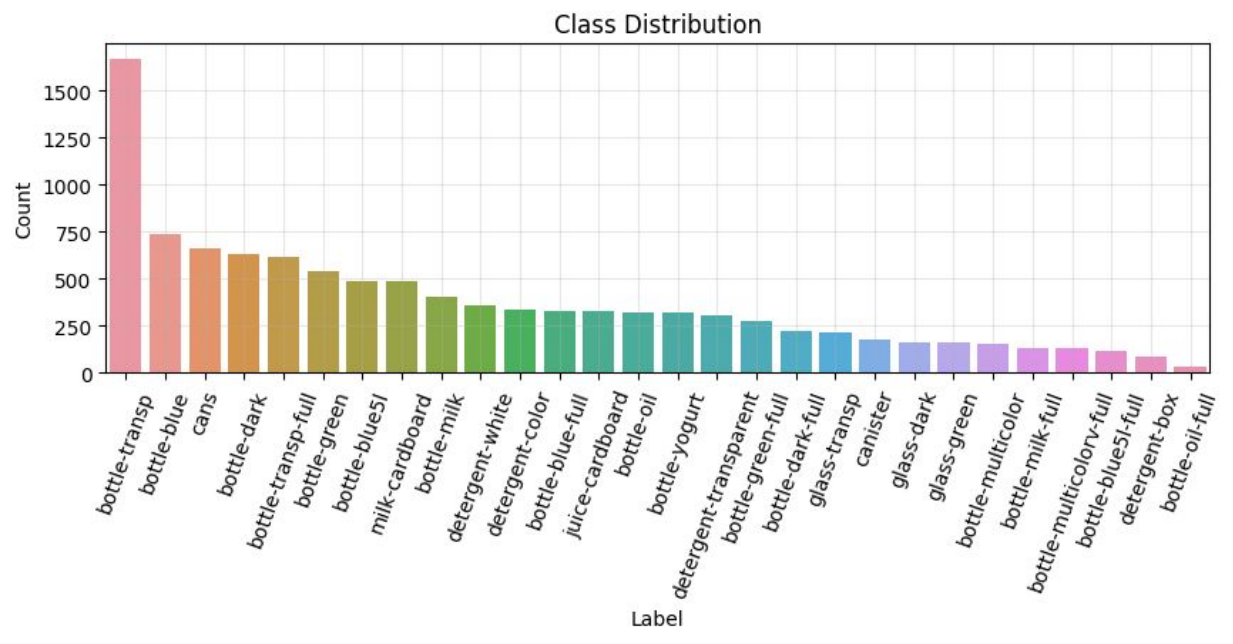

Dataset Exploration

Parent Class Distribution

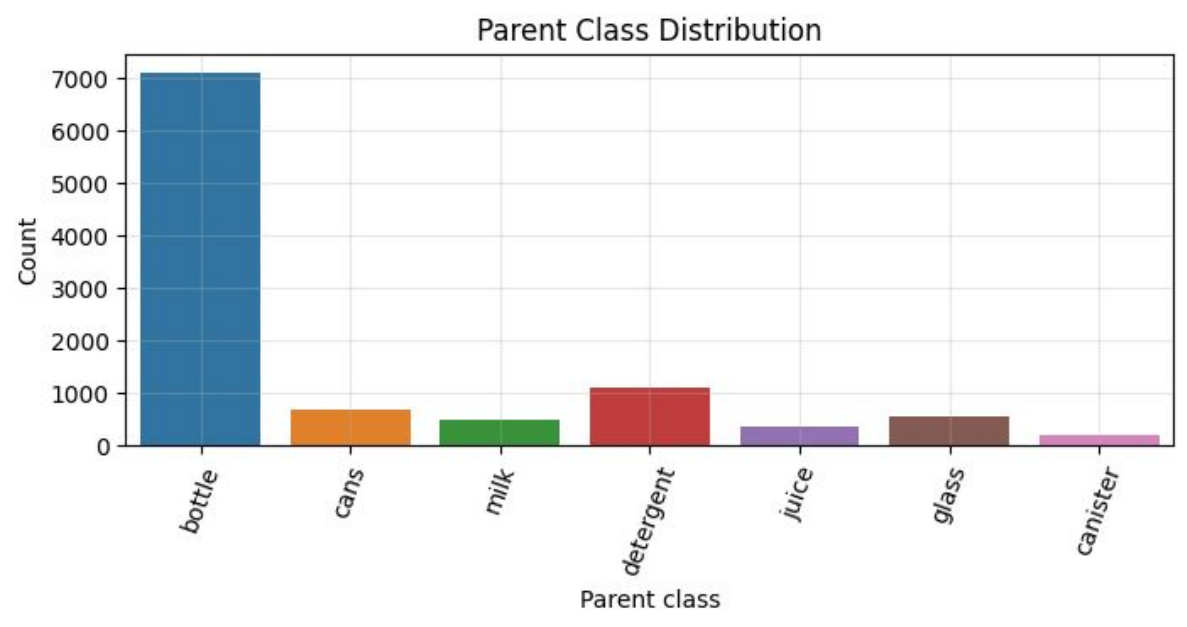

Dataset Preprocessing

Balancing over populated class (bottle-transp)

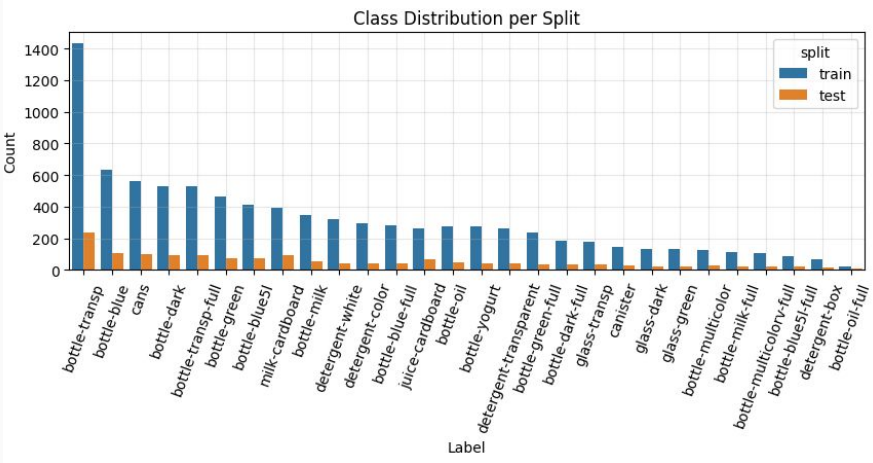

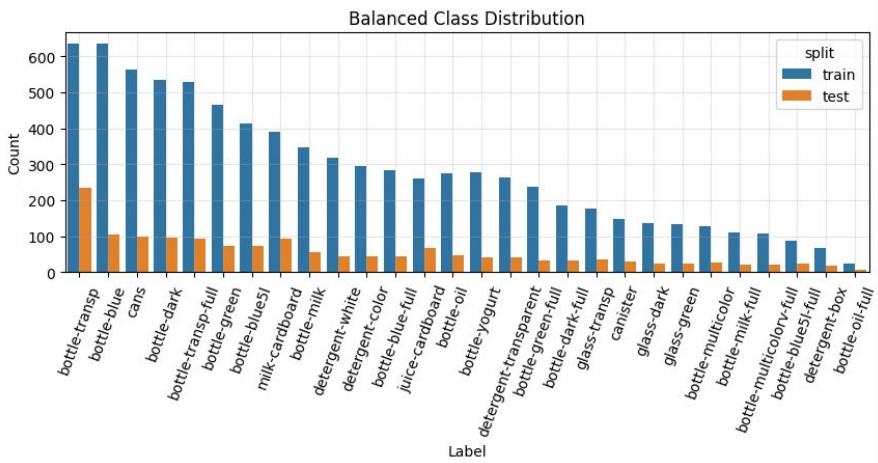

Dataset Exploration

Class Distribution Training, Test

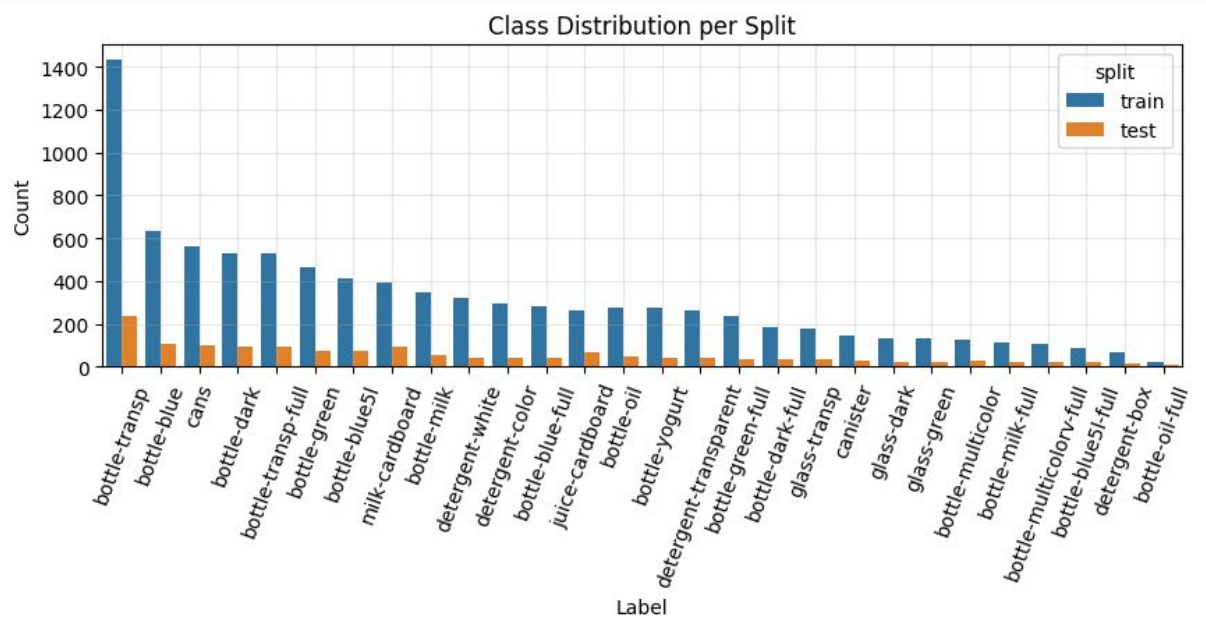

Dataset Preprocessing

Training, test, validation split (75/15/10)

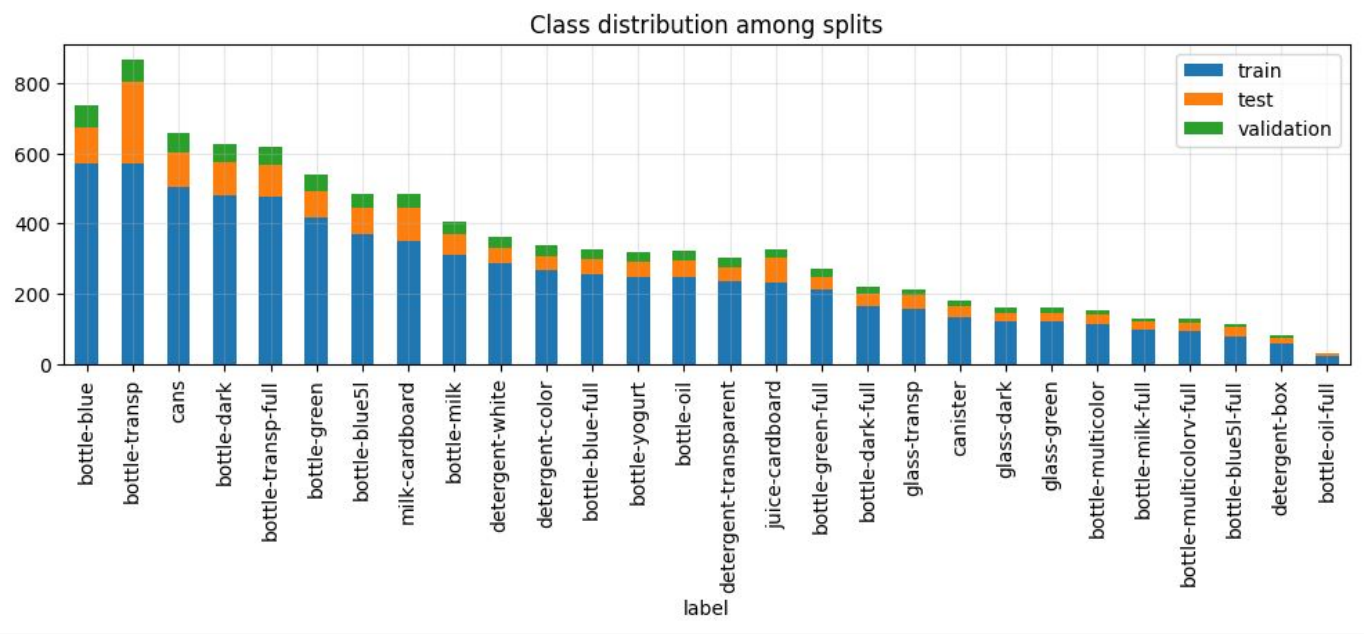

Dataset Preprocessing

Color normalization

Computed mean and standard deviation for each channel

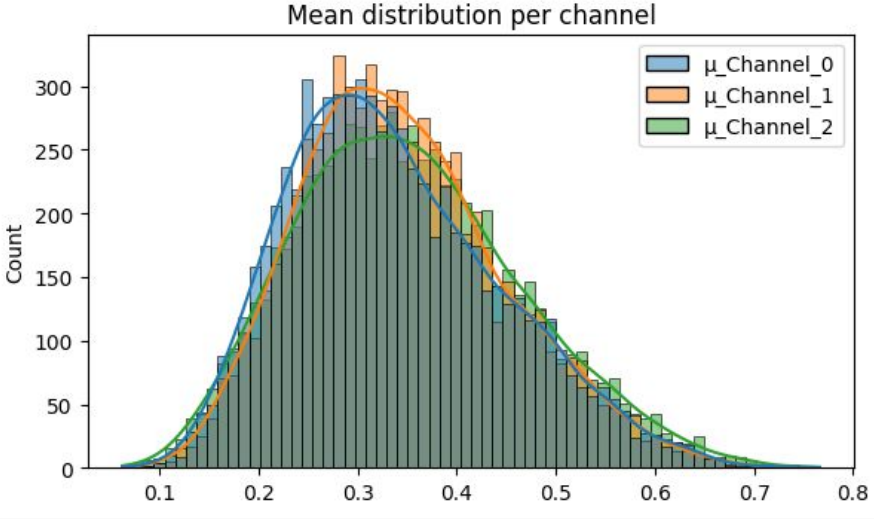

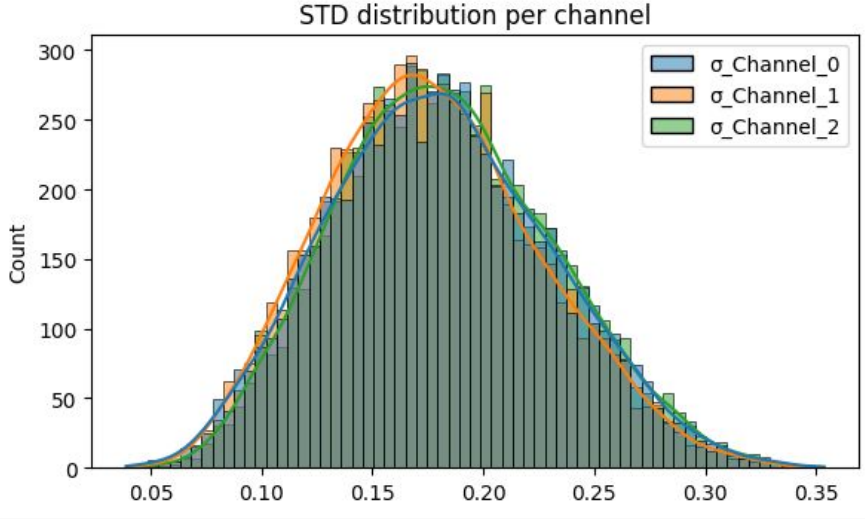

Averaged values for normalization

	Channel 1	Channel 2	Channel 3
mean	0.33	0.34	0.35
std	0.18	0.17	0.18

Dataset Preprocessing

Color normalization

Averaged values for normalization:

	Channel 1	Channel 2	Channel 3
mean	0.33	0.34	0.35
std	0.18	0.17	0.18

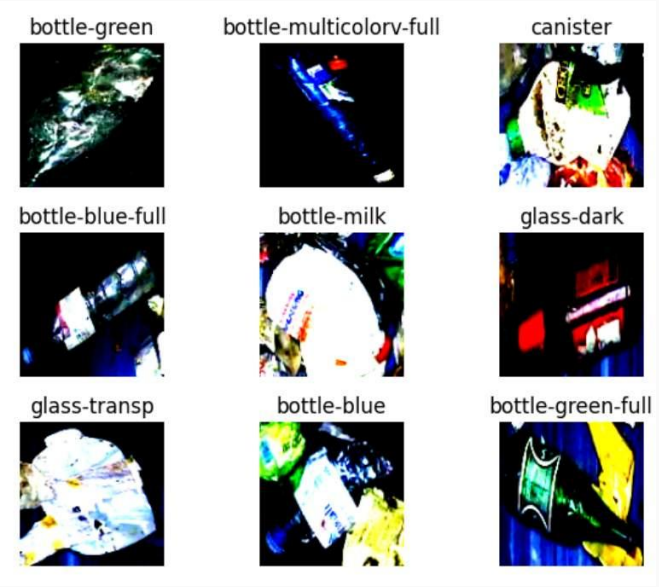

Dataset Preprocessing

Dataset augmentation

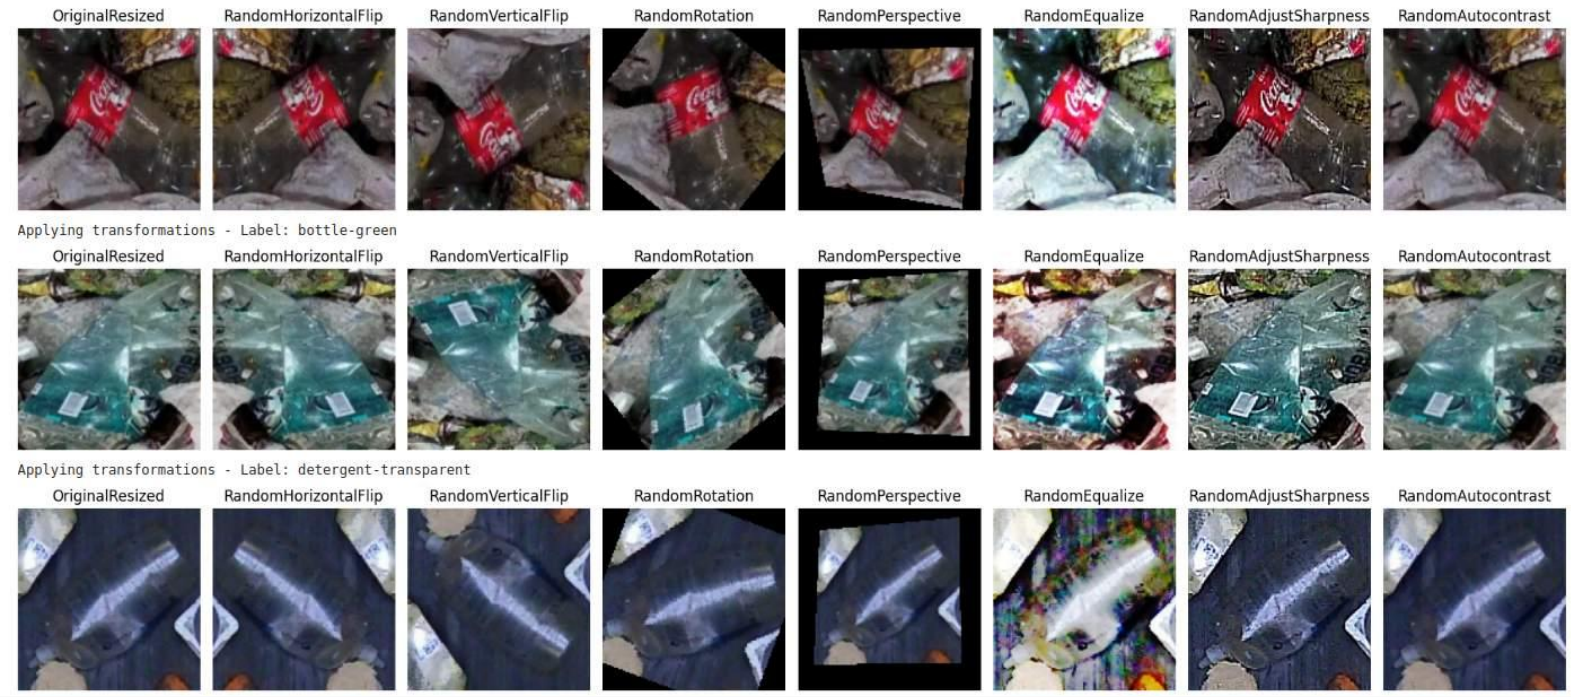

Model

CNN Architecture

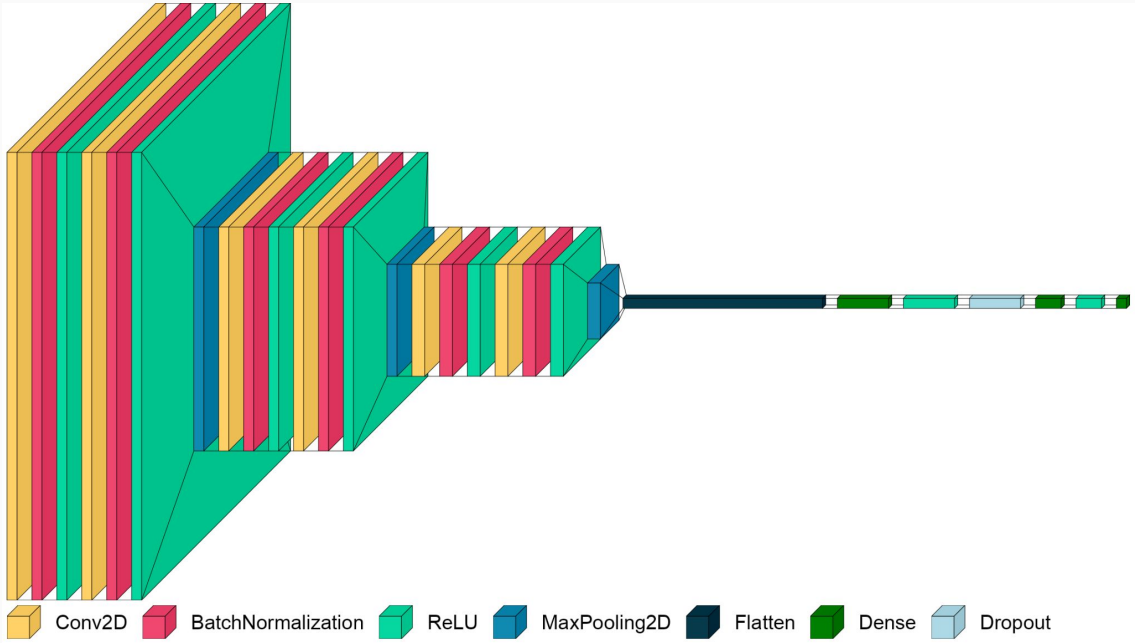

Model

Feature extraction layers

- Convolutional layers
- Batch normalization
- ReLU activation functions
- Max-pooling layers
- Dense layers

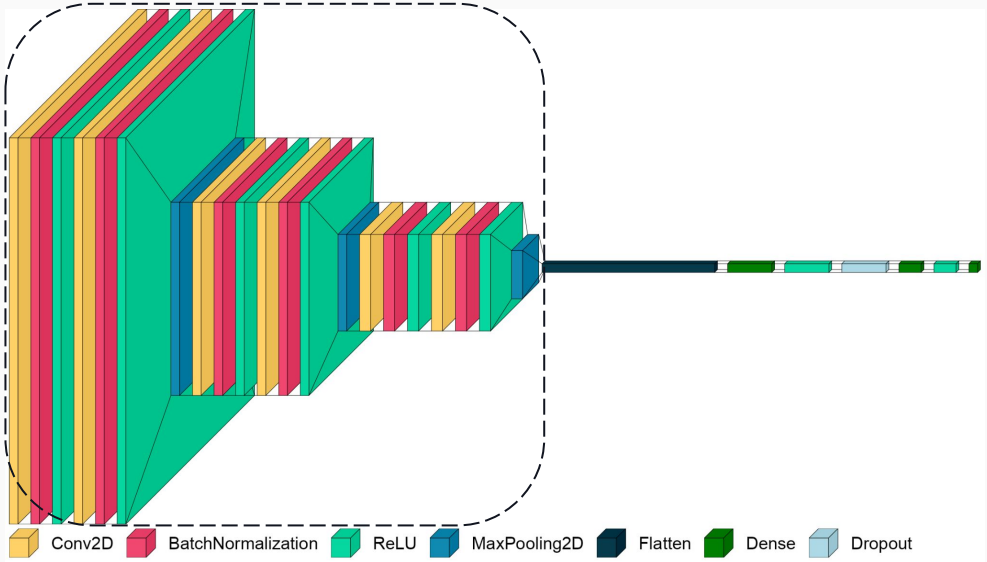

Model

Classification layers

- Flatten
- Dropout
- Dense layers
- Final Dense layers

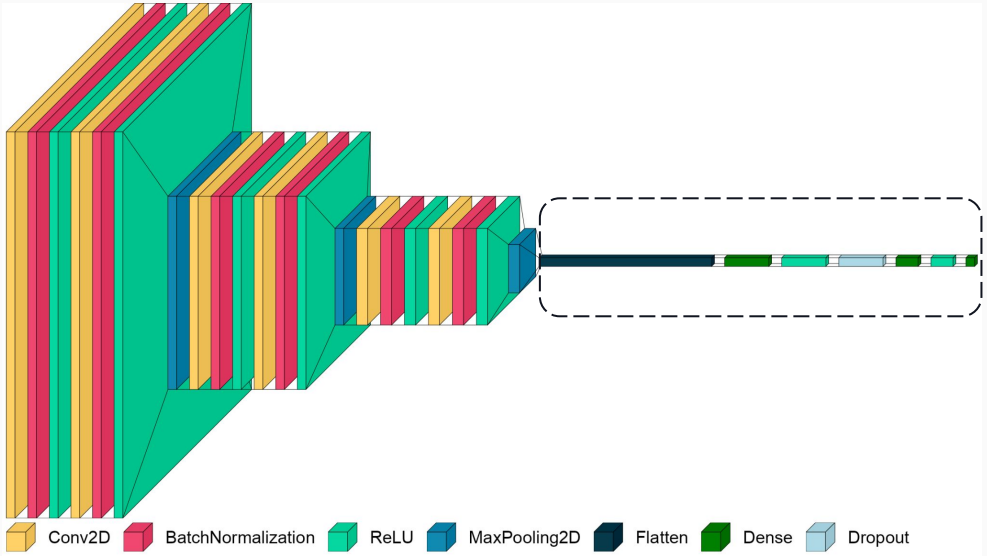

Training

Hyperparameter tuning & Early Stopping

Learning rate	0.0005, 0.0001
Convolutional unit	16, 32, 64
Batch size	16, 64, 128
Dropout rate	0, 0.2
Weight decay	0.01, 0.001
Early stopping patience	15

- Memory constraints: combinations with convolutional units 64 and batch size 128 were excluded from the grid search
- Total hyperparameter combinations is 56

Training

Hyperparameter tuning & Early Stopping

Loss and metrics on validation set for every parameter combination

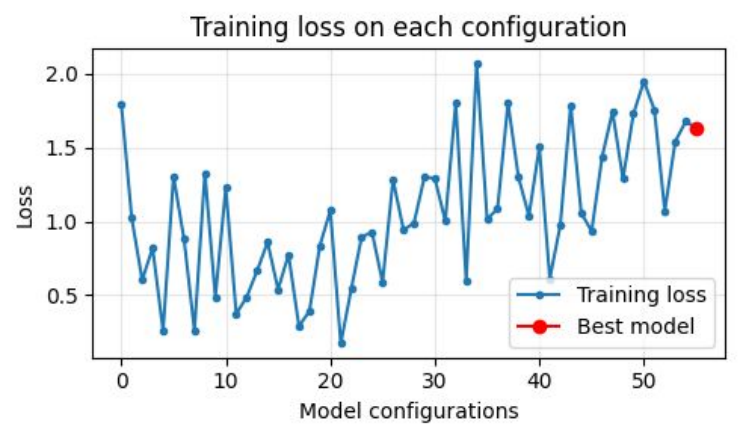

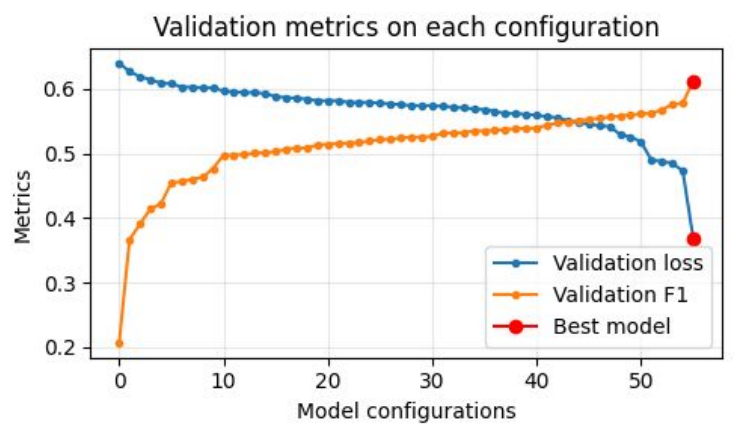

Training

Best model elected

Optimal model chosen on the highest F1-score achieved

Parameters	
Learning rate	0.0001
Convolutional unit	32
Batch size	16
Dropout rate	0
Weight decay	0.01

Metrics	
Training loss	1.798
Validation loss	0.640
Validation precision	0.641
Validation recall	0.602
Validation f1-score	0.612

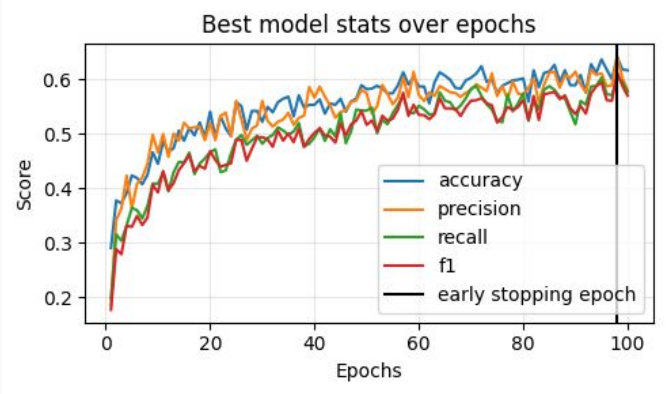

Performances

Test set classification report

	precision	recall	f1-score	support
bottle-blue	0.55	0.68	0.61	104
bottle-blue-full	0.62	0.37	0.46	43
bottle-blue5l	0.67	0.64	0.65	72
bottle-blue5l-full	0.78	0.29	0.42	24
bottle-dark	0.67	0.94	0.78	95
bottle-dark-full	0.76	0.56	0.64	34
bottle-green	0.64	0.78	0.71	74
bottle-green-full	0.92	0.65	0.76	34
bottle-milk	0.39	0.42	0.41	57
bottle-milk-full	0.82	0.43	0.56	21
bottle-multicolor	0.60	0.32	0.42	28
bottle-multicolorv-full	0.70	0.67	0.68	21
bottle-oil	0.34	0.50	0.40	48
bottle-oil-full	0.00	0.00	0.00	8
bottle-transp	0.60	0.60	0.60	234
bottle-transp-full	0.71	0.50	0.59	92
bottle-yogurt	0.36	0.33	0.35	42
canister	0.71	0.33	0.45	30
cans	0.89	0.69	0.78	98
detergent-box	0.75	0.35	0.48	17
detergent-color	0.47	0.51	0.49	43
detergent-transparent	0.35	0.22	0.27	41
detergent-white	0.31	0.51	0.39	43
glass-dark	0.94	0.60	0.73	25
glass-green	0.74	0.68	0.71	25
glass-transp	0.54	0.61	0.57	36
juice-cardboard	0.46	0.72	0.56	68
milk-cardboard	0.62	0.67	0.64	94
accuracy			0.59	1551
macro avg	0.60	0.52	0.54	1551
weighted avg	0.61	0.59	0.58	1551

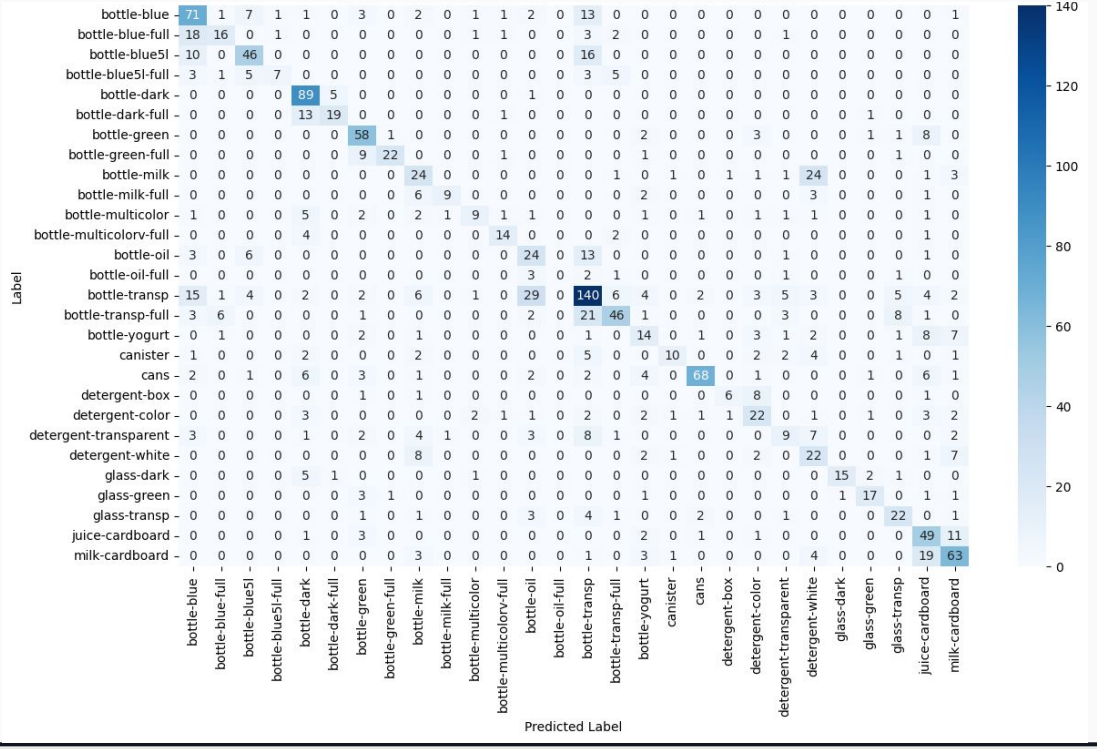

Performances

Model performance @ top-k

For this study *k=5* was chosen

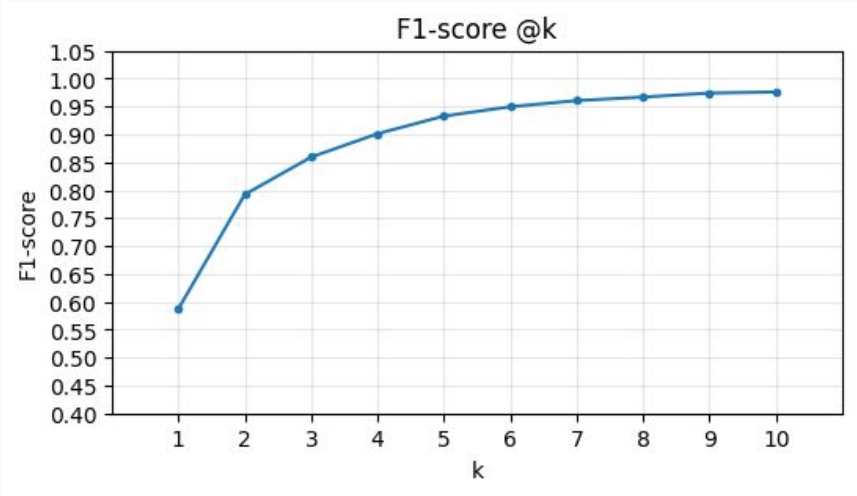

Performances

Model performance @ top-5

	precision	recall	f1-score	support
bottle-blue	0.99	0.95	0.97	104
bottle-blue-full	0.88	0.84	0.86	43
bottle-blue5l	0.91	0.97	0.94	72
bottle-blue5l-full	0.92	0.96	0.94	24
bottle-dark	0.99	0.99	0.99	95
bottle-dark-full	0.97	0.97	0.97	34
bottle-green	0.99	0.97	0.98	74
bottle-green-full	0.94	0.97	0.96	34
bottle-milk	0.98	0.96	0.97	57
bottle-milk-full	0.90	0.90	0.90	21
bottle-multicolor	0.92	0.82	0.87	28
bottle-multicolorv-full	0.78	0.86	0.82	21
bottle-oil	0.93	0.90	0.91	48
bottle-oil-full	0.44	0.50	0.47	8
bottle-transp	0.98	0.97	0.97	234
bottle-transp-full	0.91	0.90	0.91	92
bottle-yogurt	0.79	0.79	0.79	42
canister	0.71	0.73	0.72	30
cans	0.92	0.92	0.92	98
detergent-box	0.67	0.94	0.78	17
detergent-color	0.98	0.95	0.96	43
detergent-transparent	0.94	0.80	0.87	41
detergent-white	0.84	0.98	0.90	43
glass-dark	0.96	1.00	0.98	25
glass-green	1.00	0.96	0.98	25
glass-transp	0.97	0.86	0.91	36
juice-cardboard	0.93	0.97	0.95	68
milk-cardboard	0.98	0.99	0.98	94
accuracy			0.93	1551
macro avg	0.90	0.90	0.90	1551
weighted avg	0.94	0.93	0.93	1551

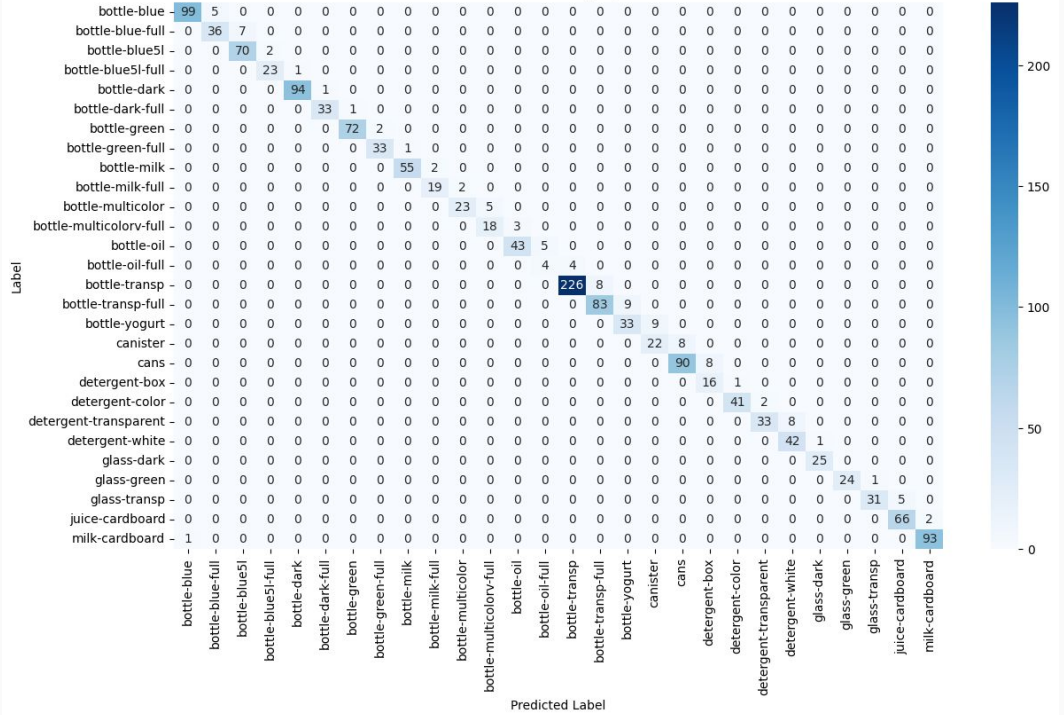

GradCAM

How to?!??!?!?

```python
def evaluate_gradcam(image, label, path):
    pred = model(image, register_hook=True)
    # label_pred = pred.argmax(dim=1)
    labels_pred = pred.topk(5)
    pred.sum().backward()

    gradients = model.get_activations_gradient()
    pooled_gradients = torch.mean(gradients, dim=[0, 2, 3])
    activations = model.get_activations(image).detach()

    for i in range(64):
        activations[:, i, :, :] *= pooled_gradients[i]

    heatmap_act = torch.mean(activations, dim=1).squeeze()
    heatmap_act = np.maximum(heatmap_act.cpu(), 0)
    heatmap_act /= torch.max(heatmap_act)

    image_rgb = cv2.imread(path)

    heatmap = cv2.resize(heatmap_act.numpy(), (image_rgb.shape[1], image_rgb.shape[0]))
    heatmap = np.uint8(255 * heatmap)
    heatmap = cv2.applyColorMap(heatmap, cv2.COLORMAP_JET)

    superimposed_img = heatmap * 0.5 + image_rgb
    superimposed_img = np.uint8(superimposed_img)
    superimposed_img = cv2.cvtColor(superimposed_img, cv2.COLOR_BGR2RGB)

    return image_rgb, heatmap_act, superimposed_img, labels_pred
```

GradCAM

Correct predictions

True label: *bottle-blue* - **Predicted:** *bottle-blue*

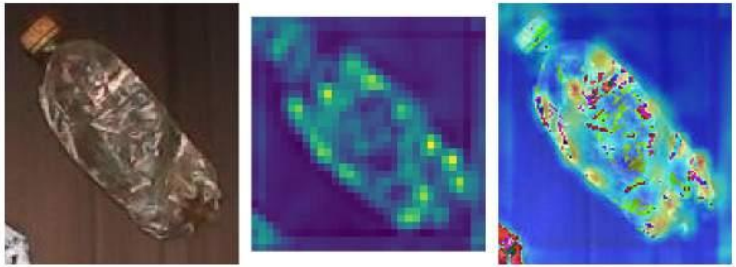

Class	Probability
bottle-blue	0.992
bottle-blue-5l	0.003
bottle-transp	0.002
bottle-oil	0.002
bottle-green	0.001

True label: *canister* - **Predicted:** *canister*

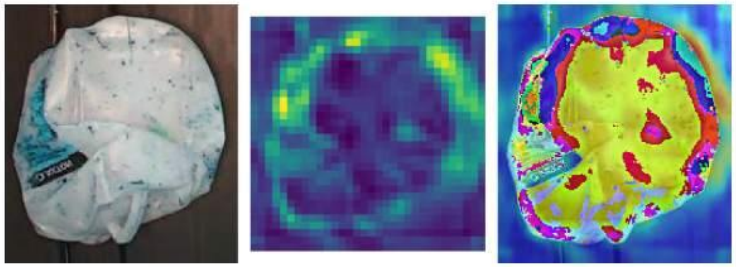

Class	Probability
canister	0.805
detergent-white	0.146
bottle-milk	0.045
bottle-multicolor	0.003
detergent-transparent	0.001

GradCAM

Incorrect predictions

True label: *bottle-blue5l-full* - **Predicted:** *bottle-blue-full*

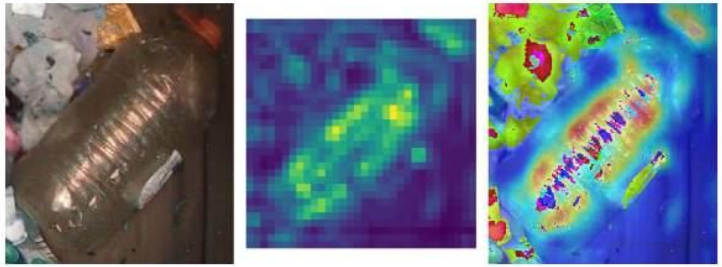

Class	Probability
bottle-blue-full	0.625
bottle-blue5l-full	0.368
bottle-blue	0.004
bottle-transp-full	0.003
bottle-green-full	0.001

True label: *milk-cardboard* - **Predicted:** *juice-cardboard*

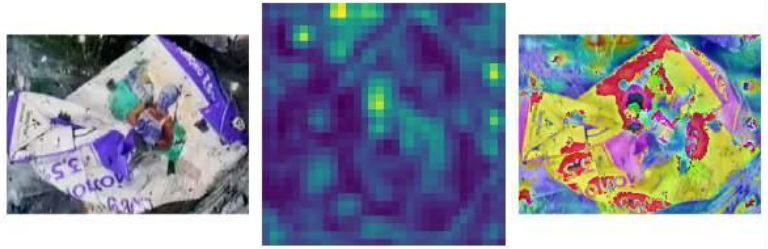

Class	Probability
juice-cardboard	0.601
milk-cardboard	0.362
detergent-white	0.037
bottle-milk	0.001
bottle-green	¡0.001

GradCAM on Video

Realtime GradCAM on Video

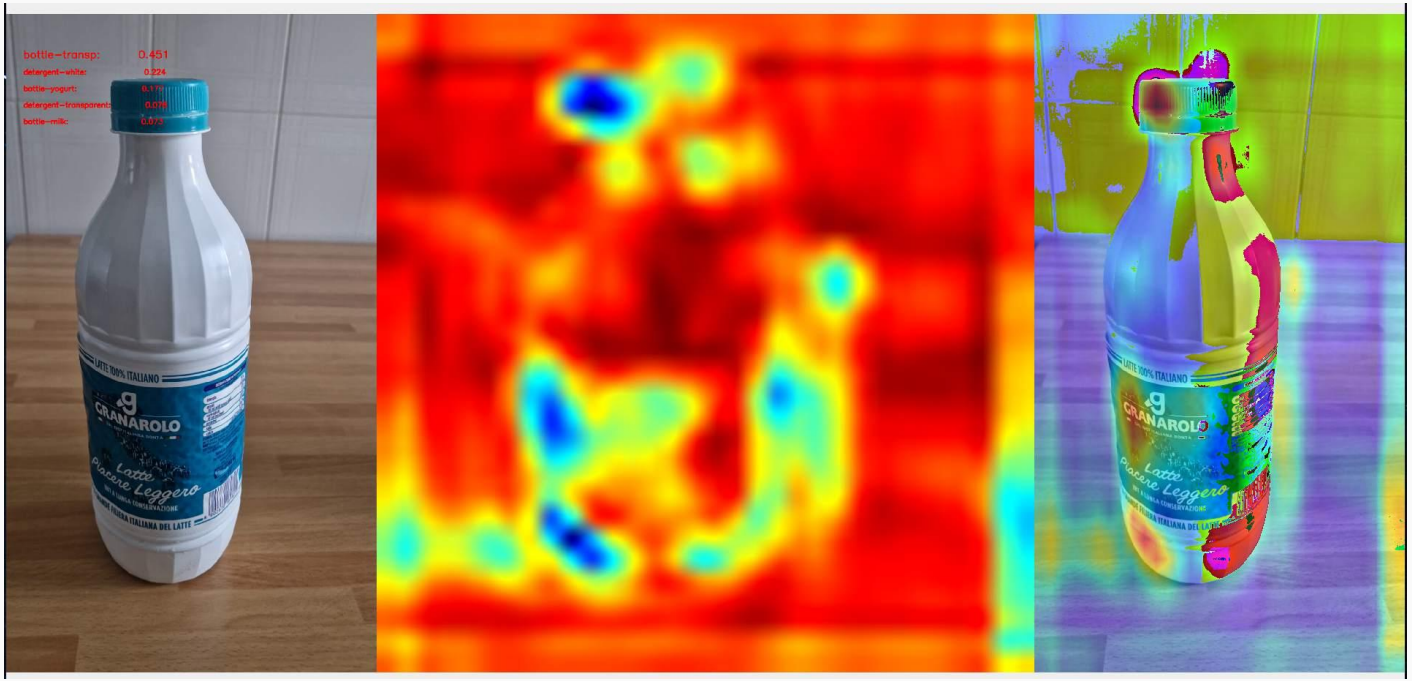

Grazie